

August 27, 2026
Hiring Manager
Confidential

Dear Hiring Manager,

I am writing to express my interest in the Senior Data Engineer role focused on Databricks Lakehouse development. With 7+ years of hands-on experience designing and delivering enterprise-scale data pipelines, Medallion Architecture, and distributed data systems on Databricks, I am confident in my ability to contribute immediately to your data engineering team.

Databricks Lakehouse & Medallion Architecture

At Ferguson Enterprises and Catalyzt Nutrition, I designed and built Databricks Medallion Architecture pipelines processing enterprise-scale ecommerce, supply chain, and operational datasets — implementing Bronze ingestion layers, Silver transformation and validation layers, and Gold dimensional models for analytics and reporting. I deployed Delta Live Tables for declarative pipeline orchestration, Unity Catalog for data governance, and Delta Lake for ACID-compliant, versioned storage across multi-cloud environments on Azure, GCP, and AWS. In my freelance practice, I extended this expertise across three client organizations: at callcathai.online, I built Databricks Medallion Architecture for call analytics with Delta Live Tables, data quality validations, and production alerting; at primelayertrades.com, I implemented Unity Catalog governance, Delta Lake storage, and CI/CD pipelines for Git-based deployment across environments; and at hereraspreadsheets.com, I developed reusable Databricks framework components, Spark optimization patterns, and streaming data pipelines feeding real-time executive dashboards.

PySpark, Spark SQL & Performance Engineering

I am hands-on with PySpark and Spark SQL across all positions — writing advanced transformations, optimizing Spark jobs for partitioning and cluster efficiency, and designing reusable framework components that scale across production environments. I have engineered dimensional data models serving CFO, VP of Finance, and EVP of Sales at organizations with large, complex datasets, and I bring strong CI/CD and Git-based deployment discipline to every deliverable. My multi-cloud experience across AWS, Azure, and GCP gives me the flexibility to operate in any Lakehouse environment your team runs on.

I would welcome the opportunity to bring this depth of Databricks expertise to your data engineering organization. I look forward to discussing how my background aligns with your platform and team.

Sincerely,
Alexander Tobin