

(757) 289-1204 | atobin@alum.utk.edu | Virginia Beach, VA

August 31, 2026
Hiring Manager

Dear Hiring Manager,

I am writing to express my strong interest in the Senior Data Engineer role. With 7+ years of hands-on Databricks experience building enterprise-scale Lakehouse pipelines, Medallion Architecture (Bronze/Silver/Gold), Delta Live Tables, Unity Catalog, and data quality monitoring frameworks, I am confident in my ability to contribute immediately to your data engineering team.

Databricks Lakehouse & Medallion Architecture

I have designed and delivered Medallion Architecture Lakehouse solutions across every position in my career. At Ferguson Enterprises, I built and migrated enterprise data infrastructure to Databricks on Azure and GCP — implementing Bronze ingestion, Silver transformation, and Gold dimensional layers for a large-scale ecommerce and supply chain operation with PySpark and Spark SQL transformations at scale. At Catalyzt Nutrition, I led Lakehouse delivery on GCP, implementing Delta Live Tables for declarative orchestration, Delta Lake for ACID-compliant versioned storage, and automated data quality validations across all pipeline layers. In my freelance practice, I extended this work across three client organizations: at callcathai.online, building Medallion pipelines with data quality monitoring and production alerting on Databricks; at primelayertrades.com, building reusable PySpark frameworks with Unity Catalog governance and CI/CD deployment; and at hereraspreadsheets.com, developing Spark optimization patterns, dimensional data models, and streaming pipelines feeding real-time executive dashboards.

PySpark, Spark SQL & Performance Engineering

I write PySpark and Spark SQL daily — designing transformations, optimizing jobs for partitioning and cluster efficiency, and engineering reusable framework components that scale across multi-node production environments. I bring strong SQL optimization, dimensional data modeling experience, and multi-cloud depth across AWS, Azure, and GCP. I have also worked with Databricks AI Genie and AI-assisted development tooling, and I maintain a strong CI/CD and Git-based deployment discipline across all deliverables.

I would welcome the opportunity to contribute to your Lakehouse engineering team and discuss how my depth of Databricks experience aligns with your needs.

Sincerely,
Alexander Tobin